

Data Mining


Lab 02


Names:

- 1) Khalil Ismail Khalil (23)
- 2) Ahmed Mohamed EL-Bawab (08)


1)Visualization:


```
data1 = dataframe.loc[1:100,2]  
Class1 = dataframe.loc[1:100,1]  
plt.scatter(data1,Class1)  
plt.show()
```






```
# Create boxplot
plt.boxplot(attributesList)
plt.show()
```






```
#Visualize correlationMatrix using imshow
plt.imshow(correlationMatrix)
```

```
<matplotlib.image.AxesImage at 0x7f49604fd588>
```




```
tmp = dataframe.loc[dataframe[1] == 'M',:]
```

```
plt.hist(tmp)
plt.show()
```



```
tmp = dataframe.loc[dataframe[1] == 'B',:]

plt.hist(tmp)
plt.show()
```



2)Dataset Splitting:

398

=====

```
[['11.69' '24.44' '76.37' ... '0.1308' '0.2803' '0.0997\n']
 ['12.23' '19.56' '78.54' ... '0.108' '0.2668' '0.08174\n']
 ['10.8' '9.71' '68.77' ... '0.04603' '0.209' '0.07699\n']
 ...
 ['12.19' '13.29' '79.08' ... '0.08187' '0.3469' '0.09241\n']
 ['6.981' '13.43' '43.79' ... '0' '0.2932' '0.09382\n']
 ['12.06' '18.9' '76.66' ... '0.05093' '0.288' '0.08083\n']]
```

171

=====

```
[['27.22' '21.87' '182.1' ... '0.2688' '0.2856' '0.08082\n']
 ['12.96' '18.29' '84.18' ... '0.06608' '0.3207' '0.07247\n']
 ['11.67' '20.02' '75.21' ... '0.0812' '0.3206' '0.0895\n']
 ...
 ['15.27' '12.91' '98.17' ... '0.1035' '0.232' '0.07474\n']
 ['14.4' '26.99' '92.25' ... '0.05563' '0.2345' '0.06464\n']
 ['12.2' '15.21' '78.01' ... '0.05556' '0.2661' '0.07961\n']]
```

```

=====
[ 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'M' 'B'
  'B' 'B' 'B' 'B' 'B' 'B' 'M' 'B' 'M' 'M' 'M' 'B' 'B' 'M' 'M' 'B' 'M' 'M'
  'B' 'M' 'M' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'B'
  'B' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'M' 'B' 'B' 'M' 'B'
  'B' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'M' 'M' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B'
  'B' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B'
  'B' 'M' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B'
  'B' 'B' 'M' 'M' 'M' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'M' 'B' 'B' 'M' 'B'
  'B' 'M' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'M' 'B' 'M' 'M' 'B' 'M' 'M'
  'M' 'B' 'B' 'B' 'M' 'M' 'M' 'M' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B'
  'M' 'B' 'M' 'B' 'B' 'B' 'M' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'M'
  'B' 'M' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'M' 'M' 'M' 'M'
  'M' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B'
  'B' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'B'
  'M' 'B' 'M' 'M' 'M' 'M' 'M' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B'
  'B' 'B' 'B' 'M' 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'M'
  'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'M' 'M' 'M' 'B' 'B' 'M' 'B'
  'B' 'M' 'B' 'B' 'M' 'M' 'B' 'M' 'B' 'B' 'M' 'B' 'B' 'M' 'B' 'B' 'M' 'M'
  'M' 'B' 'M' 'M' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'B'
  'B' 'B']
-----

```

```



=====
[ 'M' 'B' 'B' 'M' 'B' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B'
  'B' 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'M' 'M' 'B' 'M' 'B' 'M' 'M'
  'B' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'M' 'M' 'B' 'B' 'B'
  'B' 'M' 'M' 'B' 'M' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'M'
  'M' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'M' 'B' 'B'
  'B' 'B' 'M' 'M' 'M' 'B' 'B' 'B' 'B' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'B' 'B'
  'B' 'M' 'B' 'M' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'M' 'M'
  'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M'
  'M' 'B' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'M' 'B' 'B'
  'M' 'B' 'M' 'M' 'B' 'M' 'B' 'B' 'B' 'B' ]

```

3)Preprocessing:



1-train_data normalization:

```
[[2.39579978e-02 1.91521396e-02 2.60819693e-02 ... 2.23513382e-02
 3.76943844e-02 2.14313989e-02]
[5.01014308e-02 3.06340292e-02 2.34491056e-02 ... 2.43689618e-02
 7.25161987e-02 3.35892970e-02]
[1.56582095e-01 1.23021730e-01 1.66106841e-01 ... 1.45041192e-01
 2.36447084e-01 1.36255992e-01]
...
[2.56258955e-04 1.63945093e-04 1.06063345e-04 ... 1.42567114e-04
 0.00000000e+00 8.56314181e-05]
[5.62803522e-04 4.12693261e-04 4.99712704e-04 ... 6.28685995e-04
 1.58315335e-03 5.07016346e-04]
[1.92489483e-04 1.22810793e-04 1.80846337e-04 ... 1.61899616e-04
 5.06587473e-04 1.38777786e-04]]
```



2-test_data normalization:

```
[[8.46308341e-03 2.08336495e-02 2.11903117e-02 ... 1.63700773e-02
 1.95999797e-02 2.09191874e-02]
 [6.79952478e-03 2.94042092e-02 3.63556595e-02 ... 1.38397844e-02
 3.67386027e-02 2.60812706e-02]
 [5.66223285e-02 1.35354337e-01 1.36592252e-01 ... 1.05251976e-01
 1.25576294e-01 1.33781875e-01]
 ...
 [8.27279935e-05 1.00420268e-04 1.42625102e-04 ... 1.09192860e-04
 7.31529246e-05 9.17427361e-05]
 [8.79518786e-05 5.09845394e-04 5.77425613e-04 ... 2.46965165e-04
 3.16646609e-04 4.52814157e-04]
 [2.42764511e-05 1.10695291e-04 1.57699639e-04 ... 7.83575960e-05
 8.54181345e-05 1.32987952e-04]]
```



4)Classification:

1-classification models and evaluation:

a)Decision Tree:

```
from sklearn import metrics

clf = DecisionTreeClassifier()

clf = clf.fit(data_train,Class_train)

Class_pred = clf.predict(data_test)

#print(clf.get_depth)

print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
```

Accuracy: 0.9298245614035088

```
from sklearn.tree import DecisionTreeClassifier
from sklearn import metrics

clf = DecisionTreeClassifier(max_depth=5)

clf = clf.fit(data_train,Class_train)

Class_pred = clf.predict(data_test)

print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
#print(metrics.confusion_matrix(Class_test, Class_pred))
#print(metrics.f1_score(Class_test, Class_pred))
```

Accuracy: 0.9707602339181286

```
clf = DecisionTreeClassifier()

clf = clf.fit(data_train,Class_train)

Class_pred = clf.predict(data_test)

|
#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
```

```
[[94 13]
 [ 6 58]]
```



```

from sklearn import metrics

clf = DecisionTreeClassifier(max_depth=5)

clf = clf.fit(data_train,Class_train)

Class_pred = clf.predict(data_test)

#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
#print(metrics.f1_score(Class_test, Class_pred))

[[101  6]
 [ 6 58]]

```

```

clf = DecisionTreeClassifier(max_depth=5)

clf = clf.fit(data_train,Class_train)

Class_pred = clf.predict(data_test)

#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
#print(metrics.confusion_matrix(Class_test, Class_pred))
print(metrics.classification_report(Class_test, Class_pred,target_names=['M','B']))

```

	precision	recall	f1-score	support
M	0.95	0.98	0.96	107
B	0.97	0.91	0.94	64
accuracy			0.95	171
macro avg	0.96	0.94	0.95	171
weighted avg	0.95	0.95	0.95	171

b)AdaBoost:

```

# Create adaboost classifier object
abc = AdaBoostClassifier(n_estimators=50,learning_rate=1)
# Train Adaboost Classifier
model = abc.fit(data_train, Class_train)

#Predict the response for test dataset
Class_pred = model.predict(data_test)

# Model Accuracy, how often is the classifier correct?
print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))

```

Accuracy: 0.9649122807017544

```

abc = AdaBoostClassifier(n_estimators=50,learning_rate=1)
# Train Adaboost Classifier
model = abc.fit(data_train, Class_train)

#Predict the response for test dataset
Class_pred = model.predict(data_test)

# Model Accuracy, how often is the classifier correct?
#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))

[[104  3]
 [ 3 61]]

```

```

from sklearn.ensemble import AdaBoostClassifier
# Create adaboost classifier object
abc = AdaBoostClassifier(n_estimators=50,learning_rate=1)
# Train Adaboost Classifier
model = abc.fit(data_train, Class_train)

#Predict the response for test dataset
Class_pred = model.predict(data_test)

# Model Accuracy, how often is the classifier correct?
#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
print(metrics.classification_report(Class_test, Class_pred,target_names=['M','B']))

[[104  3]
 [ 3 61]]

```

	precision	recall	f1-score	support
M	0.97	0.97	0.97	107
B	0.95	0.95	0.95	64
accuracy			0.96	171
macro avg	0.96	0.96	0.96	171
weighted avg	0.96	0.96	0.96	171

- **base_estimator:** It is a weak learner used to train the model. It uses DecisionTreeClassifier as default weak learner for training purpose. You can also specify different machine learning algorithms.
- **n_estimators:** Number of weak learners to train iteratively.
- **learning_rate:** It contributes to the weights of weak learners. It uses 1 as a default value.

c)Random Forest:

```
from sklearn.ensemble import RandomForestClassifier

clf=RandomForestClassifier(n_estimators=100)

clf.fit(data_train, Class_train)

Class_pred=clf.predict(data_test)

print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
```

Accuracy: 0.9707602339181286

```
clf=RandomForestClassifier(n_estimators=100)

clf.fit(data_train, Class_train)

Class_pred=clf.predict(data_test)

#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
```

```
[[105  2]
 [ 4 60]]
```

```
clf=RandomForestClassifier(n_estimators=100)

clf.fit(data_train, Class_train)

Class_pred=clf.predict(data_test)

#print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
print(metrics.classification_report(Class_test, Class_pred,target_names=['M','B']))
```

```
[[107  0]
 [ 4 60]]
```

	precision	recall	f1-score	support
M	0.96	1.00	0.98	107
B	1.00	0.94	0.97	64
accuracy			0.98	171
macro avg	0.98	0.97	0.97	171
weighted avg	0.98	0.98	0.98	171

```

from sklearn.ensemble import RandomForestClassifier

clf=RandomForestClassifier(n_estimators=100,max_depth=5)

clf.fit(data_train, Class_train)

Class_pred=clf.predict(data_test)

print("Accuracy:",metrics.accuracy_score(Class_test, Class_pred))
print(metrics.confusion_matrix(Class_test, Class_pred))
print(metrics.classification_report(Class_test, Class_pred,target_names=['M','B']))

```

Accuracy: 0.9707602339181286

[[106 1]

[4 60]]

	precision	recall	f1-score	support
M	0.96	0.99	0.98	107
B	0.98	0.94	0.96	64
accuracy			0.97	171
macro avg	0.97	0.96	0.97	171
weighted avg	0.97	0.97	0.97	171

* some equations:

$$\text{Precision} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

$$\text{Recall} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

$$F1 = 2 \times \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

1-Hyper-parameter tuning:

```
DTC_params = {
    'max_depth' : [4,5,6,7,8,9,10]
}

ABC_params = {
    'n_estimators': [10, 50, 80, 100]
}

RFC_params = {
    'n_estimators': [10, 50, 80, 100],
    'max_depth' : [4,5,6,7,8,9,10]
}

CV_DTC = GridSearchCV(estimator=DTC_clf, param_grid=DTC_params, cv= 10)
CV_DTC.fit(data_train, Class_train)
print("DTC_best_parameters",CV_DTC.best_params_ )

CV_ABC = GridSearchCV(estimator=ABC_clf, param_grid=ABC_params, cv= 10)
CV_ABC.fit(data_train, Class_train)
print("ABC_best_parameters",CV_ABC.best_params_ )

CV_RFC = GridSearchCV(estimator=RFC_clf, param_grid=RFC_params, cv= 10)
CV_RFC.fit(data_train, Class_train)
print("RFC_best_parameters",CV_RFC.best_params_ )

DTC_best_parameters {'max_depth': 10}
ABC_best_parameters {'n_estimators': 100}
RFC_best_parameters {'max_depth': 6, 'n_estimators': 100}
```